

Best practices for handling Black Friday traffic

Tobiasz Heller
Ingrid

Agenda

- Black week traffic explanation
- Preparing load tests suite
- Finding bottlenecks in your system

E-commerce and Black Week

- +/- 10 days of very high traffic
- E-commerce merchants prepare marketing plan since summer
- New marketing strategies (influencers)

Ingrid checkout

- January 2020 - 50 RPS
- Load tests - 3000 RPS
- Black Friday 2019 - 200 RPS

Ingrid architecture


- +/- 80 microservices
- Almost all written in Go
- K8S on GCP GKE
- gRPC + google Pub/Sub for internal comm
- REST API via gRPC gateway

Load/traffic types

High stable (no spikes) traffic



Traffic spikes



Simulating high traffic

- Prepare test scheme
- Bootstrap tests on stage
- Real tests on production
- Observability!

Observability

- Prometheus + Thanos
- Metrics system need to scale with load
- Sharding of metrics

Alerting

- **Separate oncall (immediate action required) and other alerts**
- **Alerts need to be easily configurable and adjustable**

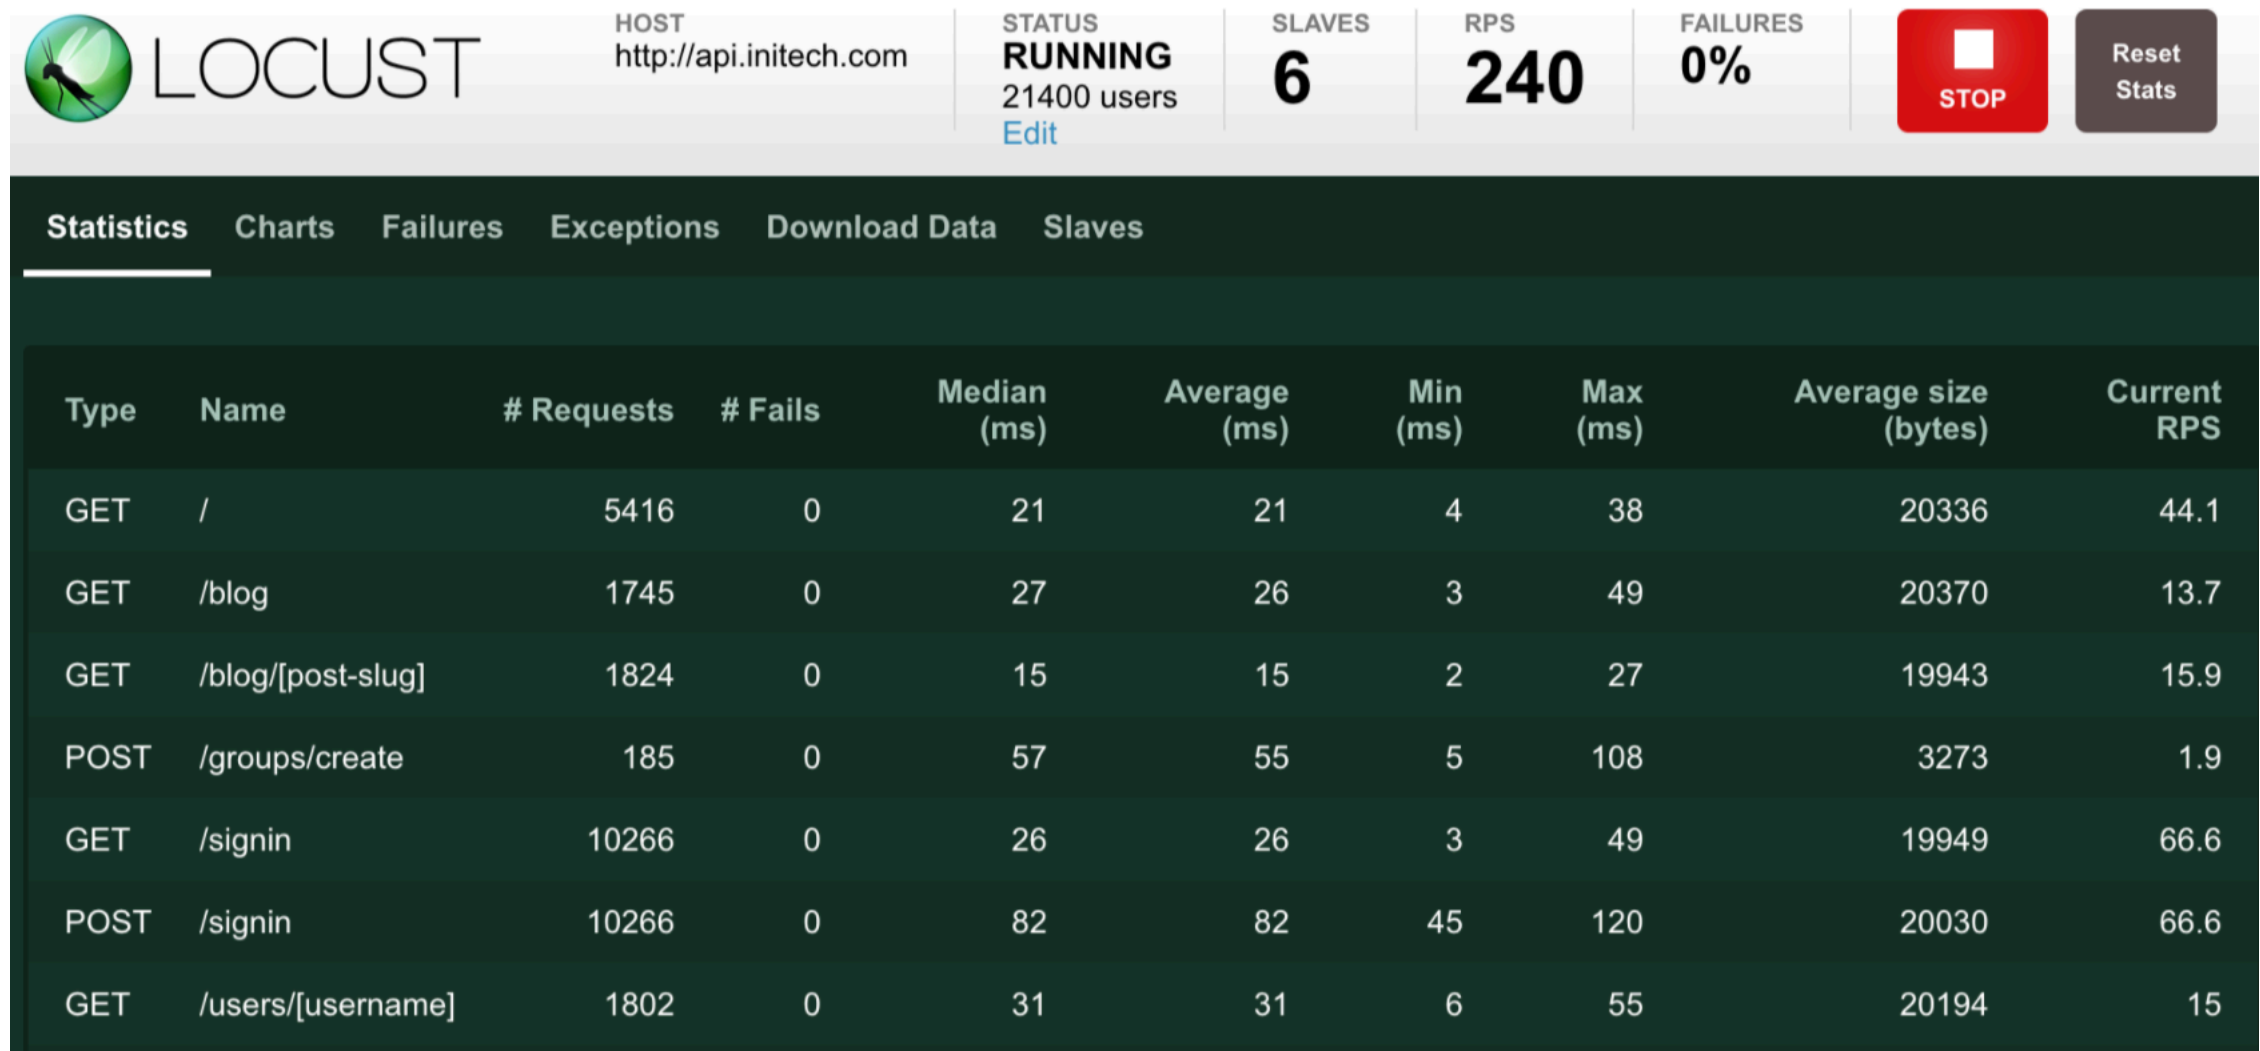
Simulating high traffic

locust.io

- Open source load testing tool
- Define user behaviour in python code
- Distributed & scalable
 - Helm chart with locust cluster
- Web-based UI

Simulating high traffic

locust.io



K8s in action

- Horizontal pod autoscaling
- Node autoscaling
 - Scaling from 10 to 100 nodes in few minutes
- Resources placeholder (pod priority)

K8s pod CPU/mem resources

- Requests resources - scheduled on a node that can give pod given resources
 - Min 100m of CPU for Go apps (on GCP)
- Limits resources - above limit value pod will be killed
 - Need to be a 20-40% bigger than requests in order to handle spike before autoscaling new pod

DB setup

- Mostly MySQL
- Some PostgreSQL
- Redis

DB bottlenecks

CPU limits

- Separate instance per service
- Monitor metrics github.com/luna-duclos/instrumentedsql
 - Review your indexes (move analytics one to other system)
- Vertical scaling

DB bottlenecks

Number of connections

- Cache connections to DB
- Separate instance per service
- Bigger pods = less amount of active connections

DB bottlenecks

GCP limits

- GCP limit number of new connections per 100s
- Bigger pods = less amount of active connections

In-memory cache in Go

- Optimised for fast reads
- Builtin sync and async refresh of entries
- Plans to open-source it

Profiling Go apps

- pprof and flame graphs are your friends!
- Json encoding & jsoniter
- Building traces span was taking most of CPU in one service
- Creating new connection to DB

Demo

github.com/tobiaszheller/talks/tree/master/black-friday/demo

Pair programming

- Less probability of errors
- Better code quality
- Also a lot of fun



Others

- Reduce components on critical path
- Async communication if possible

Future plans

- Distributed DB

Thank you