

Fault-tolerant systems in microservices world

Tobiasz Heller
Ingrid

Agenda

- What can go wrong in distributed system?
- Preventing faults on infrastructure level
- Preventing faults on application level
- Idempotent operations

**Fault tolerance is
ability of a system to continue working in
the event of partial failure.**

Wikipedia

Why do we need fault tolerance?

- Reliability
- Availability
- Security

Fault sources

- Program unrecoverable errors
- Node unavailable
- Traffic spikes
- Data or state inconsistency

O'REILLY®

Designing Data-Intensive Applications

THE BIG IDEAS BEHIND RELIABLE, SCALABLE,
AND MAINTAINABLE SYSTEMS



Martin Kleppmann

Unrecoverable errors

Unrecoverable errors

Panic

Unrecoverable errors

Panic

```
package main

func main() {
    var m map[string]float64
    m["pi"] = 3.1416
}
```

Unrecoverable errors

Panic

```
package main

func main() {
    var m map[string]float64
    m["pi"] = 3.1416
}
```



Unrecoverable errors

Panic

- Defer & recover
- Panic from a goroutine, recover in the same goroutine
 - http middleware
 - gRPC interceptors

Unrecoverable errors

OOM

Unrecoverable errors

OOM



I NEED TO KNOW WHY MOVING OUR APP TO THE CLOUD DIDN'T AUTOMATICALLY SOLVE ALL OUR PROBLEMS.



@ScottAdamsSays

Dilbert.com

YOU WOULDN'T LET ME RE-ARCHITECT THE APP TO BE CLOUD-NATIVE. JUST PUT IT IN CONTAINERS.



11-08-17 © 2017 Scott Adams, Inc./Dist. by Andrews McMeel

YOU CAN'T SOLVE A PROBLEM JUST BY SAYING TECHY THINGS. KUBERNETES.



Unrecoverable errors

OOM

- Pod resources management
 - requested
 - limits
- Multiple replicas of service



Unrecoverable errors

OOM

- Pod resources management
 - requested
 - limits
- Multiple replicas of service
- Monitoring & alerting
 - Prometheus, Grafana, Slack



Node unavailable

Node unavailable

Network faults

Network faults

https://en.wikipedia.org/wiki/Fallacies_of_distributed_computing

The fallacies [\[edit \]](#)

The fallacies are:^[1]

- The [network](#) is reliable
- [Latency](#) is zero
- [Bandwidth](#) is infinite
- The network is [secure](#)
- [Topology](#) doesn't change
- There is one [administrator](#)
- Transport cost is zero
- The network is homogeneous

Network faults

Timeouts

Network faults

Timeouts

```
func main() {  
    ctx := context.Background()  
    ctx, cancel := context.WithTimeout(ctx, 10*time.Second)  
    defer cancel()  
    req, err := http.NewRequest(http.MethodGet, "http://google.com", nil)  
    if err != nil {  
        log.Fatal(err)  
    }  
    req = req.WithContext(ctx)  
}
```

Network faults

Retries

Network faults

Retries

```
// grpc_retry "github.com/grpc-ecosystem/go-grpc-middleware/retry"
func Example_initializationWithExponentialBackoff() {
    opts := []grpc_retry.CallOption{
        grpc_retry.WithBackoff(grpc_retry.BackoffExponential(
            100 * time.Millisecond)),
    }
    grpc.Dial("myservice.example.com",
        grpc.WithUnaryInterceptor(
            grpc_retry.UnaryClientInterceptor(opts ...)),
    )
}
```

Network faults

Circuit breaker

Protects system and remotely called function by not making calls at all, when certain failure threshold is reached.

<https://martinfowler.com/bliki/CircuitBreaker.html>

Network faults

Circuit breaker

Examples in Go

github.com/rubyist/circuitbreaker

github.com/sony/gobreaker

github.com/go-kit/kit

Network faults

Service mesh

- Istio



- Linkerd



Node unavailable

Multiple service replicas

- Specify replicas of pod
- Pod anti-affinity



Event notification

- Service publishes events on action
- Low coupling
- High availability using pub/sub


Idempotent operations

Idempotent operations

No matter how many times you call the operation, the result will be the same.

Stackoverflow

Idempotent operations



Idempotent operations

Book shipment operation:

- shipment_id
- merchant_id
- transaction_id*

Enforcing retries and idempotent design

- Chaos monkey (kubemonkey)
- github.com/luno/fate
- Preemptible nodes (last max 24h)

Traffic spikes

Traffic spikes

- Horizontal pod autoscaling
- Load balancing
- Node autoscaling



State inconsistency

State inconsistency

Single SQL database

- Use transactions
- ACID

State inconsistency

Single SQL database

- (A)tomicity - all or none will occur
- (I)solation - how concurrent execution of transactions are treated (different isolation level)

Distributed databases

Eventual consistency

If no new updates are made to a given data item, eventually all accesses to that item will return the last updated value.

- Vogels, W. (2009). *"Eventually consistent"*.

Distributed database BASE

- (B)asically (A)vailable - basic reading and writing operations are available as much as possible
- (S)oft state: after some time we only have some probability of knowing the state
- (E)ventually consistent

Thank you